Qwen

Architecting Sovereignty: An Enforcement Framework for Augmented-Citizen Accounts Against Remote Control and Unauthorized Evolution

The development of advanced neurotechnologies and artificial intelligence presents profound challenges to individual autonomy and privacy [33]. This report details a comprehensive research framework designed to establish a robust, dual-layered system for securing AI-chat accounts under the paradigm of "augmented-citizen sovereignty." The core objective is to create a system where fundamental rights—such as mental privacy, identity, and autonomy—are not merely permissions to be granted but architectural impossibilities to be violated. The framework achieves this by tightly coupling a legal and governance layer, which encodes rights as legally binding covenants, with a technical enforcement layer that makes non-compliance architecturally impossible. It addresses critical threats like remote control, ghost sessions, unauthorized evolution, and malicious command insertion by embedding safeguards directly into the account's policy, its interface contracts, and the underlying protocol. This report outlines the specifications for an "Augmented-Citizen Sovereign Account," details a phased implementation roadmap for integrating this sovereignty model into AI-chat platforms, and explores the mechanisms for accountability, detection, and response.

Legal and Governance Foundations: Encoding Rights as Enforceable Covenants

The legal and governance layer of the Augmented-Citizen Sovereign Account (ACSA) framework serves as the bedrock upon which all technical enforcement is built. Its primary function is to translate abstract ethical and human rights concepts into concrete, verifiable, and legally enforceable obligations for any organization, module, or platform seeking to interact with an augmented citizen's account. This approach moves beyond traditional terms-of-service agreements, which are often one-sided and subject to unilateral change, to create a system of mutual obligation grounded in international

norms and indigenous legal traditions. The framework draws from four principal pillars: neurorights, Indigenous Data Sovereignty, Free, Prior and Informed Consent (FPIC), and EU-style high-risk AI regulations . These pillars collectively assert that neural data is distinct from other personal data due to its direct link to an individual's identity, cognitive processes, emotions, and intentions, making it more invasive and posing unique risks of behavioral profiling and manipulation [34] [69].

Neurorights provide the foundational human-centric protections, focusing on mental privacy, identity, and autonomy . The framework codifies these rights through explicit prohibitions within the PolicyProfile, such as "no neural inputs for governance," "no score-from-inner-state," and "no automated neuro-decisions" for high-impact outcomes . This directly confronts the potential for systems to make decisions based on an individual's raw neural signals without their full understanding or consent, a risk highlighted in discussions around digital medicine and AI-driven healthcare [7] [54]. By treating these prohibitions as non-negotiable invariants, the framework elevates them from best practice to a mandatory architectural constraint .

Indigenous Data Sovereignty (ID-SOV) offers a powerful political and ethical counter-narrative to corporate data ownership models [67] [68]. Grounded in principles like OCAP (Ownership, Control, Access, and Possession) from First Nations' practices and the CARE Principles for Indigenous Data Governance, ID-SOV asserts that Indigenous peoples have the right to exercise ownership over their data [16] [18] [70]. This principle is operationalized within the ACSA framework by ensuring the citizen, not the platform, is the ultimate owner and controller of their biophysical and neural data [69]. The framework requires that any neuromorphic or biosignal channel must be engaged only after obtaining explicit consent tokens, a mechanism that operationalizes the concept of dynamic, bidirectional consent rather than a one-time notice-and-consent approach [70]. This aligns with calls for Data Use Agreements (DUAs) that bind receiving entities to strict usage and security standards [70].

Free, Prior and Informed Consent (FPIC) is integrated as a mandatory guard predicate, `FpicGuard`, which acts as a gatekeeper before any sensitive operation can proceed . This elevates consent from a simple checkbox to a computationally verified, token-based requirement. Any action involving neuromorphic data or significant account-state changes must be explicitly authorized through a process that satisfies FPIC criteria, ensuring the citizen's agency is preserved. This is particularly relevant in contexts where cultural or spiritual beliefs about the head and mind, such as the Yoruba 'Ori' or Zulu/Sesotho communal wisdom, mean that interventions can be perceived as spiritually invasive, necessitating culturally appropriate consent processes [69].

Finally, the framework incorporates elements of high-risk AI regulation, such as those found in the EU's proposed AI Act, which mandate human oversight and prohibit fully automated decisions with significant impacts on individuals [10] [53]. By adopting similar "human-integrated operation" and "no automated neuro-decisions" obligations, the framework gains regulatory legitimacy and precedent, creating a strong case for why such systems should be treated as high-risk [10].

The crucial mechanism for translating these principles into an enforceable system is the HIT License and its associated SNC Invariants . The HIT License functions as the legal instrument, while the SNC Invariants are the machine-readable expressions of the legal duties. The HIT License includes provisions that make adherence to specific PolicyProfiles, such as `acs-account-v1`, a material term of the license itself. Violations, such as attempting remote control or executing an unsupervised lockout, are defined as automatic license termination events [9] [11]. This creates a powerful commercial incentive for licensees to comply, as a single violation could result in the revocation of their right to operate on augmented-citizen identities. The SNC Invariants, encoded as boolean flags within the PolicyProfile (e.g., `no_remote_control = true`), serve as the technical counterpart to the legal covenant . They are not suggestions; they are part of the specification that the licensee has agreed to honor. If a system attempts to operate under this profile while violating its invariants, it is considered "out-of-spec" . This dual approach creates a virtuous cycle: the legal framework provides the consequences for non-compliance, giving the technical rules real-world weight, while the technical enforcement provides a verifiable method for proving compliance or detecting a breach, which can then trigger legal remedies under the license.

Technical Enforcement Architecture: Proactive Prevention Through Architectural Impossibility

The technical enforcement architecture of the Augmented-Citizen Sovereign Account (ACSA) is designed to make violations of the legal and governance layer's mandates architecturally impossible. This proactive, prevention-first strategy directly addresses the user's concerns about invisible threats like malicious command insertion and ghost-user lockouts . Instead of relying on reactive measures that detect breaches after they occur, the system is constructed so that certain actions cannot be performed at all. This is achieved through a tightly integrated stack of components operating at the identity, protocol, and execution layers.

The foundation of the technical architecture is a decentralized identity anchored by a Decentralized Identifier (DID) bound to the citizen's Bostrom DID identifier . Sovereignty is asserted through a DID-bound manifest, such as the `NeuroEcoIdentityManifest`, which cryptographically declares: "Any Morpheus-Client / SNCHIT stack interacting with this account must load the specified neurorights, FPIC, and BCI 0.3 profiles, or it is out-of-spec" . This manifest acts as a public attestation of the citizen's sovereign status and the rules governing their interactions.

At the heart of the technical defense is the `PolicyProfile` object, a JSON/ALN document that serves as the "law-in-code" for the account . This is not a simple configuration file but a comprehensive, self-contained specification of rights, restrictions, and operational procedures. Its structure is critical for enforcing the framework's principles:

- **Schema References:** It contains pointers to standardized schemas like `neurorights-policy.schema.json` and `biomech-integration-policy.schema.json`, ensuring consistency and interoperability across different systems .
- **Safety Invariants:** It embeds a strict set of safety rules, including monotonicity constraints (RoH/BCI ≤ 0.3) and a prohibition on capability reduction (`no_rollbacks_or_downgrades = true`) . This ensures that the account's security envelope can only tighten over time, never weaken.
- **Required Guards:** It lists the mandatory `SafetyGuard`s that must be loaded to enforce its rules, such as `BciCeilingGuard`, `NeurorightsGuard`, `FpicGuard`, and `EcoCorridorGuard` .
- **Evolution Requirements:** It meticulously details the requirements for any state change, mandating that all persistent changes must be proposed as an `ActionProposal`, evaluated against the profile, approved by the guards, carried a DID-bound `EVOLVE` token, and logged in an audit record .

This entire architecture relies on the principle of host-local execution within a trusted enclave called the OrganicCPU, resident in the `SNCHIT / Morpheus-Client` kernel . All critical operations, especially those involving account evolution, are required to run inside this enclave, preventing direct remote access or database manipulation by external "admin" functions . This enforces the invariant of "no remote control" by design. The `is_remote_control_attempt` predicate is used to explicitly reject any evolution proposal whose origin is not host-local or properly signed by the citizen's DID, thereby thwarting attempts at hijacking through methods like clickjacking or remote commands .

The enforcement of the `PolicyProfile` is managed by a pluggable, modular system of `SafetyGuard`s . Each guard is a specialized predicate that evaluates a proposed action

against a specific subset of the policy. For example, the `BciCeilingGuard` checks if the Brain Complexity Index remains below the specified threshold, while the `NeurorightsGuard` verifies compliance with prohibitions on neural inputs for governance . Before any account state can be changed, the proposal must pass through this pipeline of guards. This pattern, where multiple independent checks are applied sequentially, is a well-established practice in secure software engineering and provides a robust defense-in-depth strategy [13].

To address the threat of "ghost sessions" or "invisible projections," the framework implements stringent controls on session management. The invariant `no_ghost_sessions = true` means that no session can be created or maintained without being written into the account's inner ledger, which is part of the tamper-evident `EvolutionAuditRecord` shard . Every session must have a visible `SessionID`, be bound to a device, and appear in the citizen's personal ledger, eliminating the possibility of hidden "shadow users" or unauthorized access channels . Similarly, the invariant `no_unsupervised_lockout = true` ensures that any lockout or removal of capabilities must either be explicitly approved via an EVOLVE token or justified by an `Errority` event, both of which are logged in the append-only ledger .

The synergy between the declarative `PolicyProfile` and the imperative Rust/ALN traits is a key design element. The trait `SovereignAccount`, implemented by any AI-chat account object, defines the exact interface through which all state changes must be mediated . Methods like `propose_account_evolution` and `apply_account_evolution` are wired directly into the guard pipeline and audit log creation logic . This type-level enforcement in Rust ensures that any code attempting to modify the account state must go through these predefined, guarded pathways, making it impossible to bypass the central security logic [77].

The Augmented-Citizen Sovereign Account Profile Specification

The Augmented-Citizen Sovereign Account (ACSA) is formally defined by a concrete `PolicyProfile` specification, which serves as the authoritative, machine-readable embodiment of the citizen's rights and the system's constraints. This profile, versioned as "v1", is a JSON/ALN object that any compliant AI-chat or neuromorphic system must load

and enforce before it can interact with the account's state . The following table details the essential fields and their functions within this profile.

Field Name	Example Value	Description
`profile_id`	`"acs-account-v1-phoenix-bostrom18"`	A unique identifier for this specific profile version, tied to the citizen's identity.
`owner_did`	`"did:bostrom:bostrom18sd2..."`	The citizen's Decentralized Identifier, binding the profile to a specific sovereign entity.
`subject_type`	`"augmented-citizen"`	Specifies the nature of the subject, triggering specialized security and governance protocols.
`applies_to`	`["ai-chat-account", "neuromorph-client"]`	A list of system types to which this profile's rules apply.
`biophysical_model_ref`	`URI to .biocompat-index-model.aln`	A reference to a model defining the biophysical compatibility constraints, specifically for RoH/BCI monotonicity.
`neurorights_policy_ref`	`neurorights-policy.schema.json`	Points to the schema that defines neuroright prohibitions (e.g., no automated neuro-decisions).
`biomech_policy_ref`	`biomech-integration-policy.schema.json`	Points to the schema defining biomechanical integration rules, such as module roles and thresholds.
`rights_floor`	`{...}`	A nested object containing the core, non-negotiable invariants.
`required_guards`	`["BciCeilingGuard", ...]`	A list of guard names that must be instantiated and loaded to enforce this profile.
`evolution_requirements`	`{...}`	A detailed specification of the process for any persistent state change.
`account_control_clauses`	`{...}`	Rules governing low-level account operations like locking and session management.
`safety_invariants`	`{...}`	Strict numerical and logical constraints on system behavior, such as ceiling values and monotonicity.

The `rights_floor` section enforces the most fundamental aspects of sovereignty. The flag `no_remote_control = true` is paramount, architecturally forbidding any actuator writes that originate from outside the host-local enclave . Similarly, `no_ghost_sessions = true` and `forbid_ghost_sessions = true` work in concert to ensure every active session is transparent and recorded, eliminating covert access channels . The clause `no_unsupervised_lockout = true` prevents administrators or automated systems from locking an account without either explicit user consent (via an EVOLVE token) or a logged justification event . Finally, `no_rollbacks_or_downgrades = true` establishes that the system's security posture can only become stronger over time; it can never be relaxed or reverted to a less secure state .

The `evolution_requirements` field dictates the precise procedure for any action that would alter the account's persistent state. This includes not just granting or revoking capabilities, but also locking the account, rotating keys, or changing privileges. Every such change must follow a multi-step process: first, it must be formulated as an `ActionProposal`; second, it must be evaluated under the current `PolicyProfile`; third, it must be approved by all active `guards` while satisfying the `safety_invariants`; fourth, it must be accompanied by a DID-bound `EVOLVE` token signed by the account owner; and finally, the successful outcome must be committed as an `EvolutionAuditRecord` in a DID-bound, append-only shard .

The `safety_invariants` section codifies the quantitative limits of the system. It specifies hard ceilings for the Brain Complexity Index (`bci_star_ceiling = 0.3`) and the Risk of Hijacking metric (`roh_ceiling = 0.3`), with a warning band established at 0.25 to allow for graceful degradation before a hard failure occurs . The flags `monotone_roh_bci = true` and `envelope_tightening_only = true` reinforce the principle of irreversible security enhancement . The `no_neural_inputs_for_governance = true` and `noscore_from_inner_state = true` flags explicitly prohibit the use of raw neural data for automated decision-making or scoring, a critical protection for mental privacy .

This comprehensive profile is then bound to the citizen's AI-chat account via a `NeuroEcoIdentityManifest` that states, in effect, that any system interacting with the account must load and honor the `acs-account-v1` profile or be considered non-compliant . This manifest acts as a cryptographic assertion, enabling other systems to verify the account's sovereign status and the rules governing it before initiating any interaction. The combination of this detailed profile with the corresponding Rust/ALN interface traits creates a complete, enforceable specification for a sovereign account.

Policy-in-Code Implementation Roadmap

Translating the conceptual framework of the Augmented-Citizen Sovereign Account (ACSA) into a functional, deployable system requires a pragmatic, phased approach focused on "policy-in-code." This roadmap outlines the necessary steps to integrate the sovereign identity and its associated protections into existing AI-chat platforms using the Morpheus/SNC ecosystem, ensuring that the legal and technical layers are developed in tandem. The goal is to create a system where the policies are not just documented but are embedded as non-bypassable constraints within the codebase itself.

Step 1: Formalize the Sovereign Account PolicyProfile. The initial phase involves codifying the `acs-account-v1` profile into a formal, machine-verifiable schema. This entails writing a `morpheus-policy-profile.schema.json` (or its ALN equivalent) that extends the base `PolicyProfile` specification with a new `account_control` section . This schema will formally define the required fields, such as `rights_floor`, `evolution_requirements`, and `safety_invariants`. Crucially, it will mandate that any `PolicyProfile` with a `subject_type` of `"augmented-citizen"` must include references to the `neurorights-policy.schema.json` and `biomech-integration-policy.schema.json` files to ensure standardization . Once the schema is defined, the canonical `acs-account-v1` profile must be published and signed under the citizen's Bostrom DID as part of a `NeuroEcoIdentityManifest` bundle, anchoring it in a verifiable piece of data .

Step 2: Define the Rust / ALN Account Interfaces. With the policy schema in place, the next step is to create the programming interface that will enforce it. This involves developing a dedicated Rust crate, tentatively named `acs-account`, within the existing Morpheus/SNC workspace . Inside this crate, a primary trait, `SovereignAccount`, must be defined. This trait will contain the signatures for all core account operations, such as `fn owner_did(&self) -> Did`, `fn policy_profile_id(&self) -> PolicyProfileId`, and `fn is_locked(&self) -> bool` . More importantly, it will define the evolution functions: `fn propose_account_evolution(...)` and `fn apply_account_evolution(...)`. These functions will be designed to run exclusively within the host-local `OrganicCPU` enclave and will serve as the sole entry point for any state-changing operation . A small adapter layer should also be created to bridge existing AI-chat session and account models to the new `SovereignAccount` trait, allowing for easier integration without requiring a complete rewrite of legacy systems . These interfaces must be designed to seamlessly integrate with the existing `SafetyGuard` pipeline and the `EvolutionAuditRecord` creation logic in `ALN.evo` .

Step 3: Wire into AI-Chat Runtimes (Host-Local Enforcement). This phase focuses on retrofitting existing AI-chat stacks to be compatible with the sovereign account model. For any chat system claiming compatibility, the implementation must require that user accounts are modeled as `SovereignAccount` instances. When an account's manifest indicates it is an augmented-citizen account (by referencing `acs-account-v1`), the system must route all high-impact account operations through the newly defined pathways . This means that actions like locking an account, changing permissions, or rotating keys must be packaged as an `EvolutionProposal` and passed through the `OrganicCPU` and its suite of `SafetyGuard`s. This effectively sandboxes these operations, preventing "admin" functions or direct database queries from bypassing the enclave and its security checks . The system must be configured to forbid any direct DB

toggling of flags or lockouts for augmented-citizen accounts, reinforcing the "no remote control" invariant .

Step 4: Implement Detection & Response for Invisible Threats. While proactive prevention is the priority, this step adds a crucial layer of verification and response for threats that might bypass initial defenses. This involves creating a conformance test suite to actively probe for vulnerabilities. Tests will be designed to attempt creating a session without logging it to verify that the `NeurorightsGuard` rejects such proposals. Other tests will fuzz `ActionProposals` to see if they can alter account state without a valid EVOLVE token or from a non-host-local origin, verifying that the `is_remote_control_attempt` predicate is effective . Additionally, hooks for an `Errority` event model must be implemented. Any observed harmful action, such as an unauthorized lockout, should trigger an `Errority` event. This event does more than just log the incident; it automatically triggers a "one-way tightening" of the system's constraints, potentially disabling noncompliant modules or profiles, thereby hardening the system against further abuse .

Step 5: Govern and Legally Anchor the Implementation. The final and most critical phase is to formally connect the technical implementation to the overarching legal and governance framework. This involves extending the existing HIT Governance Object (HGO) to include a field like `requires_policy_profile: "acs-account-v1"` for any AI-chat or neuromorphic system that operates on augmented-citizen accounts . The notes within the HIT License must be updated to state clearly that licensees are legally obligated to load and honor the `acs-account-v1` profile. Furthermore, the license should specify that violations like remote control or unsupervised lockouts constitute a material breach and lead to automatic license termination . This step transforms the system from a private security feature into a public-facing assertion of citizen sovereignty with tangible legal consequences for violators. Documentation must be provided to show how regulators, citizens, and communities can inspect the `ALN.evo` logs for account evolutions, verify that the `acs-account-v1` profile was in force at each change, and use the Errority and corridor-tightening mechanisms to revoke the operating rights of noncompliant systems .

Accountability and Response: Tamper-Evident Audit and Errority Mechanisms

While the Augmented-Citizen Sovereign Account (ACSA) framework's primary strength lies in its proactive, preventative architecture, it is fortified by a robust accountability and response layer designed to detect, log, and react to any attempted or successful violations. This layer addresses the reality that even the most secure systems can face novel or sophisticated threats. The two cornerstones of this response mechanism are the append-only, DID-bound `EvolutionAuditRecord` and the `Errority` event model, which together provide a tamper-evident trail of evidence and an automated, one-way system-hardening process.

The `EvolutionAuditRecord` is the system's permanent, cryptographically secured memory. Unlike traditional audit logs, which have been shown to be vulnerable to modification by the very systems they are meant to monitor—a common failure point in high-stakes environments like algorithmic trading [1]—the ACSA's audit records are designed for integrity. The `ALN.evo` shard where these records are stored is described as DID-bound, append-only, and cryptographically linked [2]. This means that once an event is logged, it cannot be altered or deleted. Every accepted and even denied proposal for account evolution is recorded under the citizen's DID, creating a complete and immutable chronicle of all interactions with their account . This log serves multiple purposes: it provides undeniable proof of compliance for auditors, enables regulators to investigate incidents, and gives the citizen themselves a verifiable history of all state changes, including lockouts, privilege modifications, and capability grants or revocations [27].

This tamper-evident trail is critical for detecting invisible threats. For instance, if an attacker manages to bypass initial guards and execute a "ghost session" or an unauthorized lockout, the fact that the corresponding `EvolutionAuditRecord` was created (even if the action was later rolled back) provides concrete evidence of the attempted breach. The citizen or a designated auditor can inspect the `ALN.evo` shard to see the proposal, the guards that were evaluated, the reason for rejection (if it failed), and the final outcome. This transparency allows for forensic analysis and provides the necessary evidence to trigger higher-level governance actions.

Building on this foundation of accountability is the `Errority` event model, which represents a shift from passive logging to active, automated response. In the ACSA framework, an `Errority` is not just an error message; it is a systemic event that signifies a deviation from expected safe behavior, whether caused by a bug, a malicious actor, or an unforeseen interaction. When an `Errority` is triggered—for example, by observing

an unauthorized lockout or a violation of a core invariant—it initiates a cascade of corrective actions . The most important characteristic of the `Errority` response is that it is irreversible in its effect on security posture. The system does not relax its constraints in response to a problem; it only tightens them.

For example, if an `Errority` event is raised because a particular module or profile was involved in a violation, the system's response could be to permanently disable that module or revoke the validity of that `PolicyProfile` for the citizen's account. This is a form of automated, one-way security hardening. The system becomes more restrictive and less permissive after an incident, making it harder for the same or similar attack vectors to succeed in the future. This mechanism directly supports the `no_rollbacks_or_downgrades` invariant at a meta-level; while an individual account state might be rolled back, the overall system's tolerance for risky behaviors is reduced. This aligns with the principle of "FAIR/PAIN as voluntary discipline," where the system's own internal logic punishes bad behavior by restricting the options available to the offending component .

The combination of `EvolutionAuditRecords` and `Errority` events creates a powerful feedback loop. The audit log provides the evidence, and the `Errority` model provides the consequence. This system empowers not only the citizen but also community oversight bodies or regulators who can monitor the `ALN.evo` logs. If a pattern of non-compliant behavior emerges from a particular AI-chat platform, they can use the documented evidence to enforce the terms of the HIT License, leading to the revocation of that platform's right to operate on sovereign accounts. This closes the loop between technical enforcement, legal accountability, and practical governance, transforming the framework from a theoretical construct into a living, breathing system of digital rights protection. The system is thus prepared not only for prevention but also for the detection and remediation of harm when it inevitably occurs.

Synthesis and Strategic Implications for Digital Sovereignty

The Augmented-Citizen Sovereign Account (ACSA) framework represents a holistic and ambitious synthesis of legal theory, governance policy, and advanced technical architecture. Its central contribution is the establishment of a system where fundamental rights are not negotiable permissions but are instead architecturally enforced invariants. By creating a tight coupling between legally binding covenants, embodied in the HIT

License and SNC Invariants, and non-bypassable technical enforcement, the framework moves decisively away from the reactive security paradigms of the past. It prioritizes proactive prevention, making severe violations like remote control, ghost sessions, and unauthorized rollbacks *architecturally impossible* by design [1]. This approach directly confronts the user's core concerns about invisible threats and the need for absolute sovereignty over one's digital identity and biophysical data .

The framework's strategic value lies in its dual-layered resilience. The legal layer, drawing on neurorights, Indigenous Data Sovereignty, and high-risk AI principles, provides the moral and regulatory authority, framing the protections as human rights rather than optional features [10] [38] [69]. The technical layer, with its DID-bound manifests, granular `PolicyProfiles`, host-local `OrganicCPU` execution, and a pluggable guard system, provides the verifiable mechanism to enforce these rights . The phased implementation roadmap offers a pragmatic path toward realizing this vision, starting with the formalization of the core policy and gradually integrating it into existing AI ecosystems through well-defined Rust/ALN interfaces and conformance testing . This structured approach mitigates the complexity of a large-scale overhaul and allows for incremental deployment and validation.

However, the successful long-term viability of the ACSA framework depends on addressing several critical areas. The threat model for "ghost sessions" needs further refinement to develop precise guard predicates and detection tests. The performance overhead of running complex guard evaluations in real-time for every operation must be quantified to ensure the system remains usable in latency-sensitive applications like AI-chat. Interoperability beyond the native ALN ecosystem is another key challenge; leveraging standards like Verifiable Credentials could be a promising avenue for extending trust and sovereignty claims to external services [3] [25]. Most significantly, the framework must evolve to defend against increasingly sophisticated AI-native threats, such as indirect prompt injections and protocol exploits, which represent a frontier of cyber warfare where attackers manipulate the very instructions given to an LLM [44] [62] [63]. The current design, while robust against many known attack vectors, would need specific enhancements to counter these advanced forms of malicious command insertion.

In conclusion, the ACSA framework provides a powerful and comprehensive blueprint for achieving genuine digital sovereignty in an age of autonomous systems. Its greatest strength is the symbiotic relationship between its legal and technical components, creating a system that is both ethically grounded and mechanically robust. While challenges related to scalability, interoperability, and evolving threats remain, the framework's focus on architectural impossibility sets a new standard for digital rights protection. If successfully implemented and adopted, it could serve as a foundational

model for a new generation of AI systems that prioritize human dignity, autonomy, and control above all else.

Reference

1. Building Tamper-Proof Audit Trails: What Three 2025 Trading ... https://dev.to/veritaschain/building-tamper-proof-audit-trails-what-three-2025-trading-disasters-teach-us-about-cryptographic-378g
2. Building Cryptographic Audit Trails for SEC Rule 17a-4 - Dev.to https://dev.to/veritaschain/building-cryptographic-audit-trails-for-sec-rule-17a-4-a-technical-deep-dive-4hbp
3. Yves Goeleven's Post - LinkedIn https://www.linkedin.com/posts/goeleven_verifiable-credentials-is-an-interesting-activity-7436380299539865601-TdjF
4. [XLS] Definitions - ITU https://www.itu.int/dms_pub/itu-t/oth/0a/0d/t0a0d00000a0004xlse.xls
5. Appl. Sci., Volume 13, Issue 21 (November-1 2023) – 483 articles https://www.mdpi.com/2076-3417/13/21
6. Proceedings of 10th International Conference on Coastal and ... https://link.springer.com/content/pdf/10.1007/978-981-97-5353-6.pdf
7. Untitled | PDF | Medicine | Educational Technology - Scribd https://www.scribd.com/document/636488929/Untitled
8. Sensors, Volume 21, Issue 24 (December-2 2021) – 347 articles https://www.mdpi.com/1424-8220/21/24
9. Exploring Software Licensing Issues Faced by Legal Practitioners https://dl.acm.org/doi/full/10.1145/3643766
10. Licensing high-risk artificial intelligence: Toward ex ante justification ... https://www.sciencedirect.com/science/article/pii/S0267364923001097
11. Impact of License Choice on Open Source Software Development ... https://www.researchgate.net/publication/220433105_Impact_of_License_Choice_on_Open_Source_Software_Development_Activity
12. Health Information Technologies—Academic and Commercial ... https://link.springer.com/article/10.1186/s13012-016-0495-2

13. A Survey on Modern Code Review: Progresses, Challenges and ... https://arxiv.org/html/2405.18216v1
14. Hewlett Packard Enterprise Unveils Self-Driving Network Vision https://www.linkedin.com/posts/sudheermatta_arubacentral-junipermist-discover2025-activity-7402589656262836225-Lmlb
15. Le Ton Beau de Marot in Praise of The Music of Language (Douglas ... https://www.scribd.com/document/942726048/Le-Ton-Beau-de-Marot-in-Praise-of-the-Music-of-Language-Douglas-Richard-Hofstadter-Z-Library
16. Measuring the diversity of data and metadata in digital libraries https://www.researchgate.net/publication/389215054_Measuring_the_diversity_of_data_and_metadata_in_digital_libraries
17. (PDF) Unlocking the potential of digital collections. A call to action. https://www.academia.edu/127869131/Unlocking_the_potential_of_digital_collections_A_call_to_action
18. A Framework for FAIR and CLEAR Ecological Data and Knowledge https://www.researchgate.net/publication/394458374_A_Framework_for_FAIR_and_CLEAR_Ecological_Data_and_Knowledge_Semantic_Units_for_Synthesis_and_Causal_Modelling
19. Enabling an Extensible Network of Biodiversity Data Records as ... https://www.researchgate.net/publication/362457857_Digital_Extended_Specimens_Enabling_an_Extensible_Network_of_Biodiversity_Data_Records_as_Integrated_Digital_Objects_on_the_Internet
20. (PDF) Harvestable Metadata Services Development: Analysis of Use ... https://www.researchgate.net/publication/372153932_Harvestable_Metadata_Services_Development_Analysis_of_Use_Cases_from_the_World_Data_System
21. (PDF) Umbrella Data Management Plans to Integrate FAIR Data https://www.researchgate.net/publication/374001573_Umbrella_Data_Management_Plans_to_Integrate_FAIR_Data_Lessons_From_the_ISIDORe_and_BY-COVID_Consortia_for_Pandemic_Preparedness
22. VeriTrust—A Blockchain-Based Framework for Fake News Detection https://www.mdpi.com/1999-5903/17/10/448
23. VeriTrust—A Blockchain-Based Framework for Fake News Detection https://www.researchgate.net/publication/395994803_Self-Sovereign_Identities_and_Content_Provenance_VeriTrust-A_Blockchain-Based_Framework_for_Fake_News_Detection
24. (PDF) Interoperable Architecture for Digital Identity Delegation for AI ... https://www.researchgate.net/publication/

399962473_Interoperable_Architecture_for_Digital_Identity_Delegation_for_AI_Agents_with_Blockchain_Integration

25. Verifiable Credentials Data Model v2.0 - W3C https://www.w3.org/TR/vc-data-model-2.0/
26. Verifiable Credentials JSON Schema Specification - W3C https://www.w3.org/TR/vc-json-schema/
27. (PDF) Policy-Bound Triple-Entry Receipts for Autonomous Commerce https://www.researchgate.net/publication/400806541_Policy-Bound_Triple-Entry_Receipts_for_Autonomous_Commerce
28. COAST Services: Achieving Service Customization and Policy ... https://escholarship.org/content/qt9h97w18p/qt9h97w18p.pdf?t=oeisl5
29. [XLS] Explanation of threat matrix - UNECE Wiki https://wiki.unece.org/download/attachments/44269802/TFCS-05-05-Rev2%20%28Chair%29%20Table%20on%20CS%20threats%20-%20with%20explanations%20and%20draft%20definitions.xlsx?api=v2
30. LLMOrbit: A Circular Taxonomy of Large Language Models - arXiv https://arxiv.org/html/2601.14053v1
31. [PDF] LLMOrbit - From Scaling Walls to Agentic AI Systems - arXiv.org https://arxiv.org/pdf/2601.14053
32. [PDF] Can LLMs Do Rocket Science? Exploring the Limits of Complex ... https://arxiv.org/pdf/2602.03630
33. Protecting Brain Privacy in the Age of Neurotechnology https://www.researchgate.net/publication/384971350_Protecting_Brain_Privacy_in_the_Age_of_Neurotechnology_Policy_Responses_and_Remaining_Challenges
34. Advocating for neurodata privacy and neurotechnology regulation https://www.researchgate.net/publication/373854405_Advocating_for_neurodata_privacy_and_neurotechnology_regulation?_tp=eyJjb250ZXh0Ijp7InBhZ2UiOiJzY2llbnRpZmljQ29udHJpYnV0aW9ucyIsInByZXZpb3VzUGFnZSI6bnVsbH19
35. [XLS] Explanation of threat matrix - UNECE Wiki https://wiki.unece.org/download/attachments/46792867/TFCS-07-10-Rev1%20%28Sec%29%20Mitigations%20table%20-%20finally%20revised%20during%20TFCS-07%20%28Day%202%29.xlsx?api=v2
36. [PDF] PRivacy Enabled Capability In Co-Operative Systems and Safety ... https://cordis.europa.eu/docs/projects/cnect/1/224201/080/deliverables/001-PRECIOSAD16ResearchContributionToV2XPrivacyAndRoadmapsv10.pdf

37. Indigenous data sovereignty in intangible cultural heritage governance https://www.researchgate.net/publication/392659984_Indigenous_data_sovereignty_in_intangible_cultural_heritage_governance_A_complementary_approach_to_public-private_partnerships
38. (PDF) Indigenous Data Sovereignty and Policy - Academia.edu https://www.academia.edu/44929752/Indigenous_Data_Sovereignty_and_Policy
39. (PDF) Māori Data Sovereignty and Privacy Tikanga in Technology ... https://www.researchgate.net/publication/369378933_Maori_Data_Sovereignty_and_Privacy_Tikanga_in_Technology_Discussion_Paper
40. From prompt injections to protocol exploits: Threats in LLM-powered ... https://www.sciencedirect.com/science/article/pii/S2405959525001997
41. [PDF] Stealthy Prompt Injection Attacks Targeting LLM-based Phishing ... https://arxiv.org/pdf/2602.05484
42. What Features in Prompts Jailbreak LLMs? Investigating the ... - arXiv https://arxiv.org/html/2411.03343v2
43. Unleashing the potential of prompt engineering for large language ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12191768/
44. Agentic AI Security: Threats, Defenses, Evaluation, and Open ... - arXiv https://arxiv.org/html/2510.23883v2
45. Summon a demon and bind it: A grounded theory of LLM red teaming https://pmc.ncbi.nlm.nih.gov/articles/PMC11734899/
46. StruQ: Defending Against Prompt Injection with Structured Queries https://arxiv.org/html/2402.06363v2
47. A Systematic Review of Prompt Injection Attacks on Large Language ... https://ieeexplore.ieee.org/iel8/6287639/11323511/11359704.pdf
48. Srs Document | PDF | Unit Testing | Specification (Technical Standard) https://www.scribd.com/document/95135752/Srs-Document
49. Definition Key Points Cross-references - Springer Link https://link.springer.com/content/pdf/10.1007%2F978-0-387-39940-9_919.pdf
50. Wise Tender 310510 | PDF | Arbitral Tribunal - Scribd https://www.scribd.com/document/64235700/Wise-Tender-310510
51. Digital Health Policy and Cybersecurity Regulations Regarding ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11999725/
52. (PDF) Digital Health Policy and Cybersecurity Regulations ... https://www.researchgate.net/publication/389897859_Digital_Health_Policy_and_Cybersecurity_Regulations_Regarding_Artificial_Intelligence_AI_Implementation_in_Healthcare

53. Global toolkit on AI and the rule of law for the judiciary - unesco https://unesdoc.unesco.org/ark:/48223/pf0000387331
54. Ethical and legal challenges of artificial intelligence-driven healthcare https://pmc.ncbi.nlm.nih.gov/articles/PMC7332220/
55. [PDF] Business and Human Rights: A Guide for Financial Institutions | MSCI https://www.msci.com/downloads/web/msci-com/discover-msci/events/event-assets/2026/february/Business%20and%20Human%20Rights_%20A%20Guide%20for%20Financial%20Institutionspdf%20(1)%202.pdf
56. Biometric data landscape in Southeast Asia - ScienceDirect.com https://www.sciencedirect.com/science/article/pii/S0267364924001602
57. How Prompt Injections Gradually Evolved Into a Multistep Malware ... https://arxiv.org/html/2601.09625v2
58. SoK: Prompt Hacking of Large Language Models - arXiv.org https://arxiv.org/html/2410.13901v1
59. From Prompt Injections to Protocol Exploits: Threats in LLM ... - arXiv https://arxiv.org/html/2506.23260v1
60. [PDF] A Systematic Literature Review on LLM Defenses Against Prompt ... https://arxiv.org/pdf/2601.22240
61. Threats in LLM-Powered AI Agents Workflows - arXiv.org https://arxiv.org/html/2506.23260v2
62. Adaptive Attacks on Trusted Monitors Subvert AI Control Protocols https://arxiv.org/html/2510.09462v1
63. [PDF] The Promptware Kill Chain: How Prompt Injections ... - arXiv.org https://arxiv.org/pdf/2601.09625
64. Security in Hierarchical Autonomy Evolution of AI Agents - arXiv https://arxiv.org/html/2603.07496v1
65. Multi-Agent Framework for Threat Mitigation and Resilience in AI ... https://arxiv.org/html/2512.23132
66. [PDF] Automating Agent Hijacking via Structural Template Injection https://arxiv.org/pdf/2602.16958
67. Recognising Indigenous data sovereignty and implementing ... https://www.sciencedirect.com/science/article/pii/S2949840625000300
68. Indigenous data sovereignty in intangible cultural heritage governance https://www.cambridge.org/core/services/aop-cambridge-core/content/view/5F9E115795FA06E77A05C8066D5A5D6B/S0940739125100064a.pdf/indigenous_data_sovereignty_in_intangible_cultural_heritage_governance_a_complementary_approach_to_publicprivate_partnerships.pdf

69. A Neuroethical Inquiry into Consent and Cognitive Liberty in South ... https://link.springer.com/article/10.1007/s12152-025-09615-3
70. Ethical, legal, and policy challenges in field-based neuroimaging ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11157461/
71. [PDF] UNDP Handbook obalka final:Layout 1 https://www.undp.org/sites/g/files/zskgke326/files/publications/2008_UNDP_CSO-Handbook-Public-Oversight-of-the-Security-Sector-2008.pdf
72. FortiSIEM Rules - Fortinet Document Library https://docs.fortinet.com/document/fortisiem/7.1.3/built-in-detection-rules
73. [PDF] PART 2/2 ANNEX to the REPORT FROM THE COMMISSION TO TH https://www.europarl.europa.eu/cmsdata/286238/AMPR%20-%204.COM_COM(2024)0401(ANN02PAR02)_EN.pdf
74. Exploit Tool Invocation Prompt for Tool Behavior Hijacking in LLM ... https://arxiv.org/html/2509.05755v3
75. Prompt Injection Attacks in Large Language Models and AI Agent ... https://www.mdpi.com/2078-2489/17/1/54
76. Prompt Injection Attacks and Defenses in LLMs - Rohan's Bytes https://www.rohan-paul.com/p/prompt-injection-attacks-and-defenses
77. Rust traits. Is it customary to define them in the dependency library ... https://stackoverflow.com/questions/63319764/rust-traits-is-it-customary-to-define-them-in-the-dependency-library-or-can-yo